

## Part-I: What is an Agent?

Short Answer: Text-to-Task

An LLM agent is an AI system that utilizes a LLM as its core computational engine to exhibit capabilities beyond text generation.

Exercise: Create an agent that can do some math

[Do this exercise \(18 mins\) to understand what an agent does](#)

Code: [https://github.com/rohanbalkondekar/finetune\\_llama2](https://github.com/rohanbalkondekar/finetune_llama2)

Data: [rohanbalkondekar/generate\\_json](#)

## Part-II: RAG overview

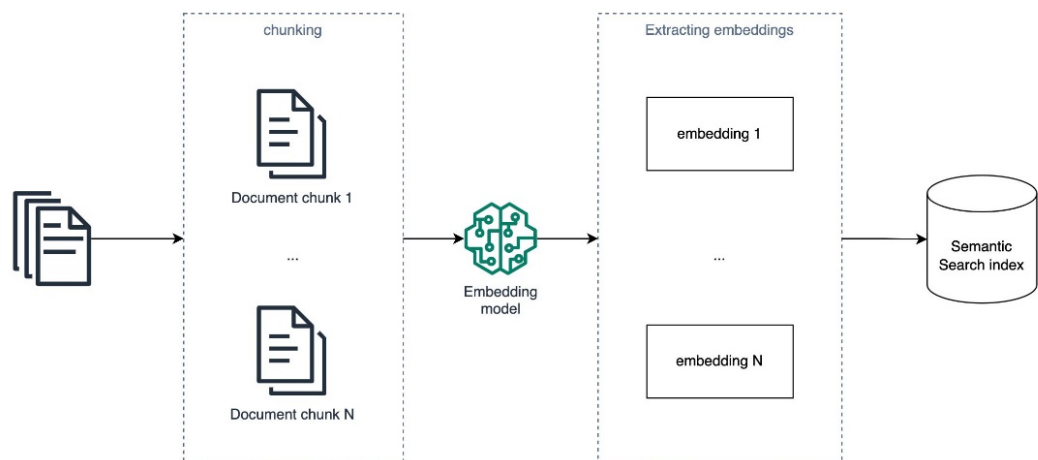
Lets review the RAG design pattern and discuss its limitations on analytical questions. Also cover a more versatile architecture that overcomes these limitations.

### Overview of RAG

RAG solutions are inspired by [representation learning](#) and [semantic search](#) ideas that have been gradually adopted in ranking problems (for example, recommendation and search) and natural language processing (NLP) tasks since 2010.

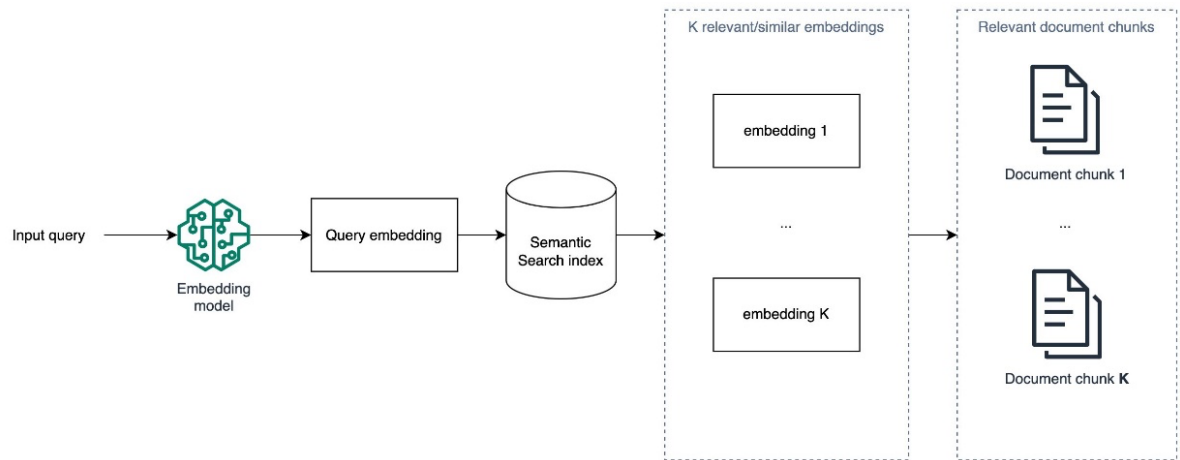
The popular approach used today is formed of three steps:

1. An offline batch processing job ingests documents from an input knowledge base, splits them into chunks, creates an embedding for each chunk to represent its semantics using a pre-trained embedding model, such as [Amazon Titan](#) embedding models, then uses these embeddings as input to create a semantic search index.

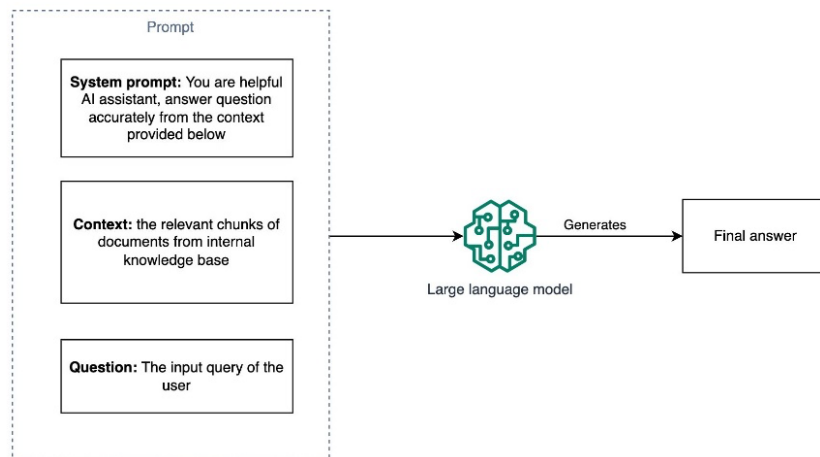


2. When answering a new question in real time, the input question is converted to an embedding, which is used to search for and extract the most similar chunks of documents using a similarity metric, such as cosine similarity, and an approximate nearest neighbors

algorithm. The search precision can also be improved with metadata filtering.



3. A prompt is constructed from the concatenation of a system message with a context that is formed of the relevant chunks of documents extracted in step 2, and the input question itself. This prompt is then presented to an LLM model to generate the final answer to the question from the context.



With the right underlying embedding model, capable of producing accurate semantic representations of the input document chunks and the input questions, and an efficient semantic search module, this solution is able to answer questions that require retrieving existent information in a database of documents. For example, if you have a service or a product, you could start by indexing its FAQ section or documentation and have an initial conversational AI tailored to your specific offering.

## PART-III: RAG LIMITATIONS on Semantic Search (Multi Document)

### 1. Multi Document

Although RAG is an essential component in modern domain-specific AI assistants and a sensible starting point for building a conversational AI around a specialized knowledge base, it can't answer questions that require:

- I. Scanning
- II. Comparing
- III. Reasoning across all documents in your knowledge base simultaneously, especially when the augmentation is based solely on semantic search.

#### A. First Check: Single document as knowledge base (complex data)

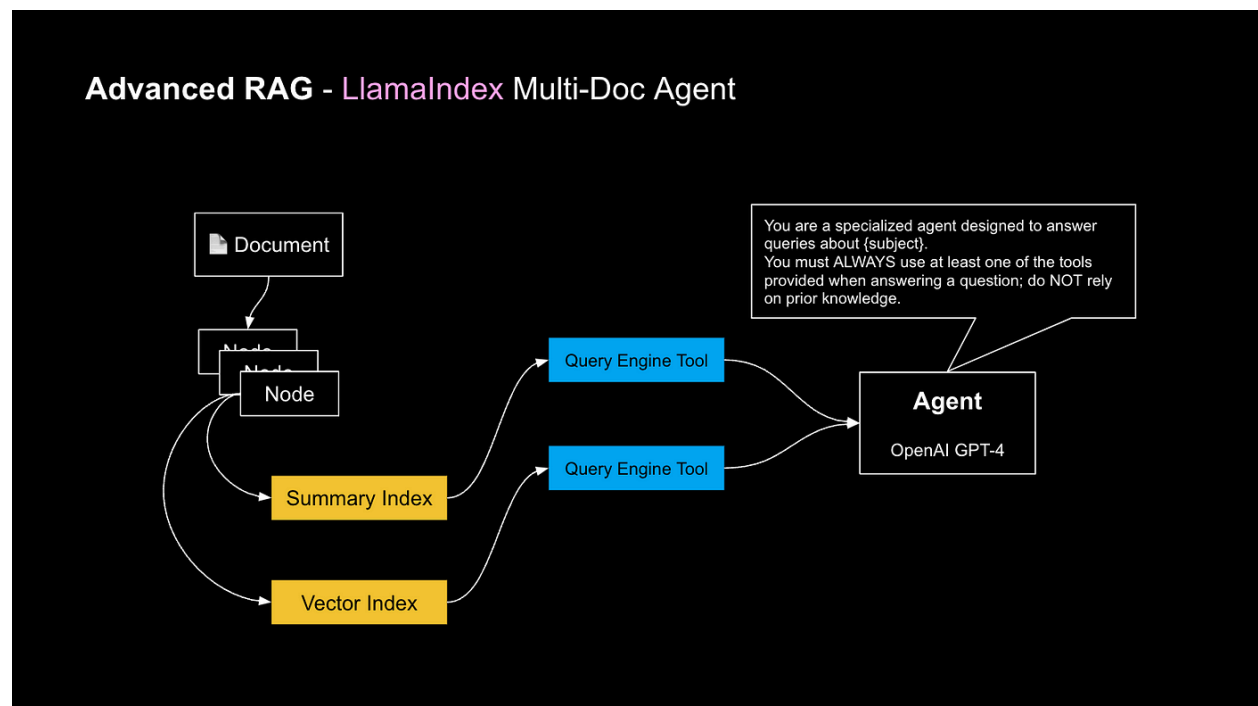
<https://ai.gopubby.com/advanced-rag-semi-structured-data-with-langchain-ce46c8baa6cf>

<https://ai.gopubby.com/advanced-rag-multi-modal-rag-with-gpt4-vision-e4c11229682c>

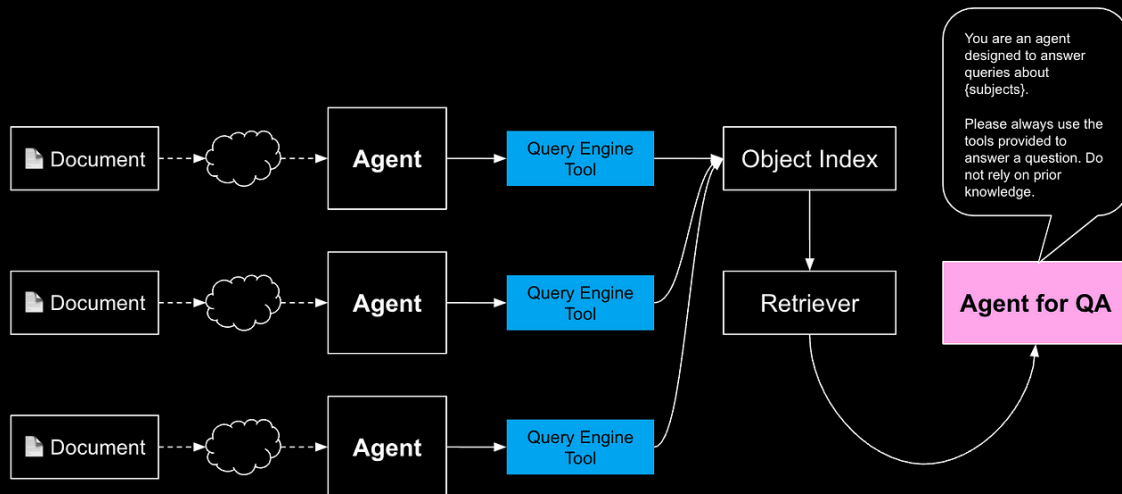
#### B. Advanced RAG – Multi Documents Agent with LlamaIndex

Sample implementation:

[https://github.com/sugarforever/Advanced-RAG/blob/main/03\\_llama\\_index\\_multi\\_doc\\_agent.ipynb](https://github.com/sugarforever/Advanced-RAG/blob/main/03_llama_index_multi_doc_agent.ipynb)



## Advanced RAG - LlamaIndex Multi-Doc Agent



- VectorStoreIndex
- SummaryIndex
- ObjectIndex
- QueryEngineTool
- OpenAIAgent
- FnRetrieverOpenAIAgent

Full Explanation

<https://blog.gopenai.com/advanced-rag-multi-documents-agent-with-llamaindex-43b604f84909>

Additional Resources from Official Documentation

[https://docs.llamaindex.ai/en/stable/examples/agent/multi\\_document\\_agents.html](https://docs.llamaindex.ai/en/stable/examples/agent/multi_document_agents.html)

In this guide, you learn towards setting up an agent that can effectively answer different types of questions over a larger set of documents.

These questions include the following

- QA over a specific doc
- QA comparing different docs
- Summaries over a specific doc
- Comparing summaries between different docs

We do this with the following architecture:

- Setup a “document agent” over each Document: each doc agent can do QA/summarization within its doc
- Setup a top-level agent over this set of document agents. Do tool retrieval and then do CoT over the set of tools to answer a question.

If you want additional features:

[https://docs.llamaindex.ai/en/stable/examples/agent/multi\\_document\\_agents-v1.html](https://docs.llamaindex.ai/en/stable/examples/agent/multi_document_agents-v1.html)

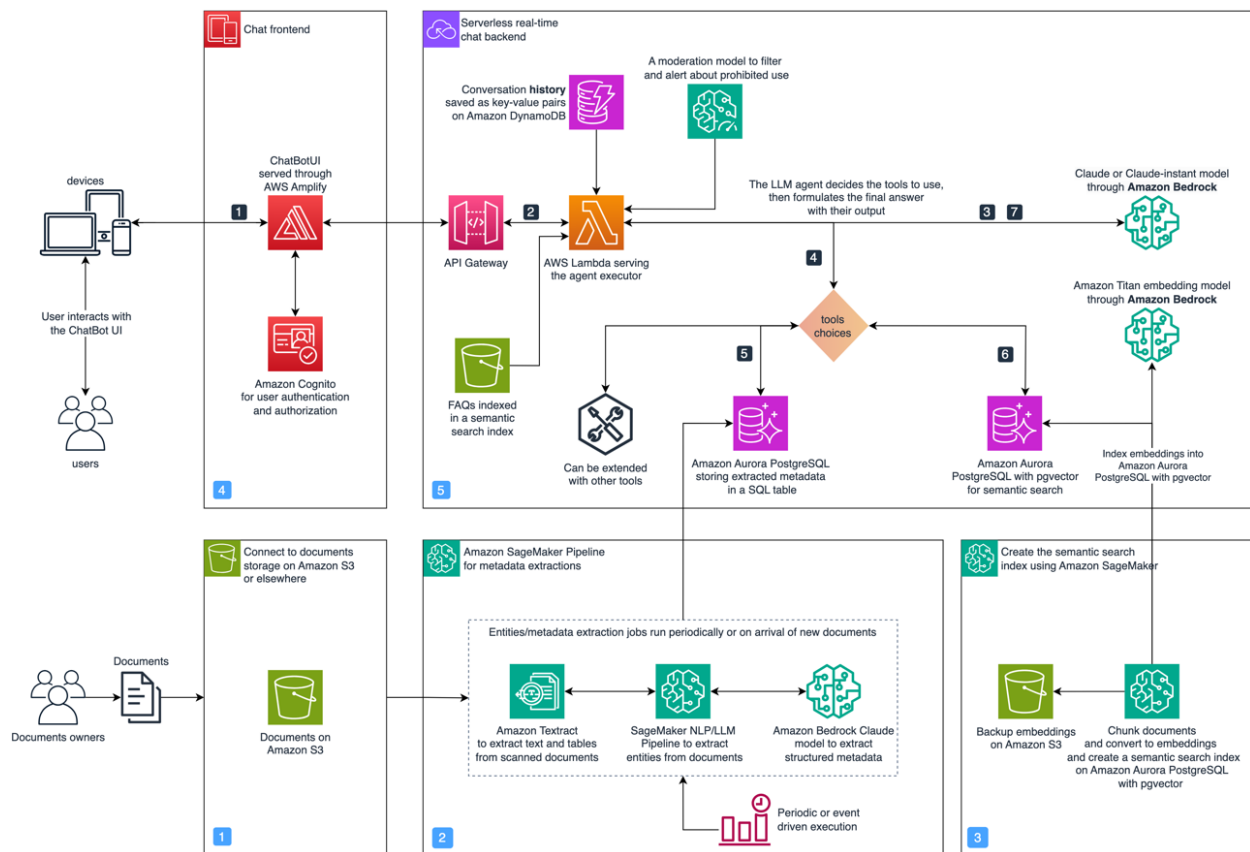
- Reranking during document (tool) retrieval
- Query planning tool that the agent can use to plan

We do this with the following architecture:

- setup a “document agent” over each Document: each doc agent can do QA/summarization within its doc
- setup a top-level agent over this set of document agents. Do tool retrieval and then do CoT over the set of tools to answer a question.

## PART-IV: RAG LIMITATIONS on Semantic Search (STRUCTURED DATA)

### Entity extraction, SQL querying, and agents with Amazon Bedrock



### Full Stack Example with Public Data